


Collaborative Filtering

PCTeam: Nicolas Durat, Christian Kayo, Marc Schachtsiek

Problem Statement

Matrix factorization

- Embedding model
- Regression-type optimization problem
- Given the rating matrix $\mathbf{A} \in \mathbb{R}^{m \times n}$, (m : number of users / n number of items) the model learn
 - User-Embedding matrix $\mathbf{U} \in \mathbb{R}^{m \times d}$, where row i is the embedding for user i
 - Item-Embedding matrix $\mathbf{V} \in \mathbb{R}^{n \times d}$, where row j is the embedding for item j
- The embeddings are learned such that the product $\mathbf{UV}^T = \hat{\mathbf{A}}$ approximates \mathbf{A}
- Objective function : (Root) Mean Square Error (RMSE)



$$\min_{U, V} \|\mathbf{A} - \mathbf{UV}^T\|_F^2$$

Approach comparison

Non-negative Matrix Factorization (NMF)

Learns parts-based representation that combine additively

Summary [1]

- Non-negative constraint on matrices
- Decomposes into basis (W) and encoding matrix (H)
- Algorithms converge to local minima and results can vary by starting point

Optimization Problem

$$\min_{W+, H+} \frac{1}{2} \|X - WH\|_F^2 + \mu \|H^T H - I\|_F^2$$

Regularization

- *Orthogonality*
 - Introduces a stronger clustering effect to NMF
 - Can be applied to basis, encoding or both matrices (ONMF)
- *Other regularization methods such as Lasso can also be applied*

Sparse Principal Component Analysis (SPCA)

Learns holistic representations that combine linearly

Summary

- Reduction dimension technique adding sparsity structures to construct PCs
- Maximizing variance along a vector and reconstruct PCA with few inputs
- (Penalized) Regression-type optimization problem using elastic-net penalty [2]

Optimization Problem

$$\operatorname{argmin}_{U, V} \|A - AUV^T\|_F^2 + \lambda \sum_{j=1}^k \|v_j\|^2 + \sum_{j=1}^k \lambda_{1,j} \|v_j\|_1$$

Regularization

- *Lasso*
 - Shrinks the coefficients towards zero
 - Select variables to produce sparse model
 - *Ridge*
 - Variable selection not limited by the n° of observations
 - Ensure PCs reconstruction
- } *Elastic Net*

Experiments & Evaluation Metrics

Non-negative Matrix Factorization (NMF)

- **Compare different NMF implementations**
 - (Accelerated) Multiplicative Update (MU)
 - (Accelerated) Hierarchical Alternating Least Squares (HALS)
 - Evaluate convergence rate
 - Evaluate computational complexity
- **Compare NMF to ONMF [3]**
 - Evaluate sparsity of encoding matrix via Misidentification Rate (MR)
 - Evaluate Root Mean Square Error (RMSE)

Sparse Principal Component Analysis (SPCA)

- **Compare different regularized SPCA**
 - Focus on SPCA (Elastic Net) and GPower (Lasso) [4]
 - Evaluate Percentage of Explained Variance (PEV)
 - Evaluate sparsity of encoding matrix via Misidentification Rate (MR)
 - Evaluate Root Mean Square Error (RMSE)

Comparison of approaches

Root Mean Square Error (RMSE)
Misidentification Rate (MR)
Computational Complexity (Time & Space)
Interpretability of U and V matrices

Implementation

Non-negative Matrix Factorization (NMF)

MU update of H

$$H^{(k)} = H^{(k)} \frac{W^{(k)T} M}{W^{(k)T} W^{(k)} H^{(k)}}$$

HALS update of H

for $p = 1:d$ **do**

$$H_{p:}^{(k)} = \max(0, H_{p:} + A_{p:} - B_{p:} \times H)$$

with $[A = W^T M, B = W^T W]$

Normal NMF

- Update W exactly once, then H exactly once

Accelerated NMF [5]

- Update W at least once, then H at least once
 - Computations for A, B only performed once per consecutive update
- Hybrid criterion to determine update count with α, ϵ parameters
 - Fixed count $l = 1: \lfloor 1 + \alpha \rho_H \rfloor$ with $\rho_H = \frac{K+md}{nd+n}$
 - Dynamic stopping $\|H^{(k,l+1)} - H^{(k,l)}\|_F \leq \epsilon \|H^{(k,1)} - H^{(k,0)}\|_F$

Adaptation for input with missing data

Sparse Principal Component Analysis (SPCA)

I. Update B (PC loading) given fixed A (PC weights)

- A. Elastic net / Lasso optimization by gradient descent
 - a. Initialization weights A , L1 & L2 penalties and Gram matrix $X^T X$
 - b. Calculate gradients
 - c. Update loadings B

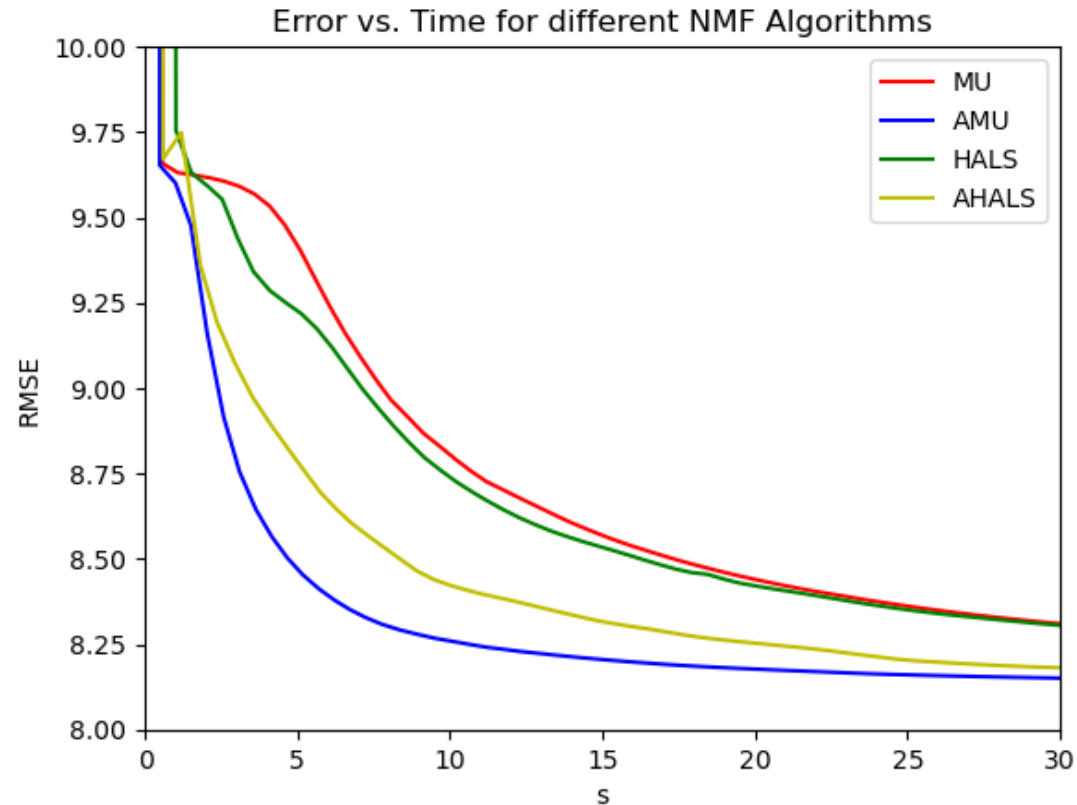
II. Update A (PC weights) given fixed B (PC loading)

- A. Compute SVD : $(X^T X)B = UDV^T$
 - 1) Update $A = UV^T$

III. Repeat Steps I. and II. until convergence

IV. Eigenvectors (PCs axes) ridge normalization

Preliminary Results



- Sparse rating matrix dataset [6]
- Generally decreasing RMSE
- Accelerated versions outperform original algorithms
- Performance of HALS not clear
 - HALS outperforms MU → as expected
 - AHALS outperforms AMU → not as expected
 - Implementation problem? Related to the dataset? Could a different metric provide better insight?
- Possible comparison metric to show relative improvement w.r.t. time

$$E(t) = \frac{e(t) - e_{\min}}{e(0) - e_{\min}} \quad \text{where} \quad e(t) = \|M - WH\|_F$$

References

- [1] Daniel D. Lee, H. Sebastian Seung; Learning the parts of objects by non-negative matrix factorization. *Nature*, 1999; 401: 788–791
- [2] Hui Zou, Trevor Hastie and Robert Tibshirani; Sparse Principal Component Analysis. *Journal of Computational and Graphical Statistics* 2015; 15 (2): 265–286
- [3] Keigo Kimura, Yuzuru Tanaka, Mineichi Kudo; A Fast Hierarchical Alternating Least Squares Algorithm for Orthogonal Nonnegative Matrix Factorization. *Proceedings of the Sixth Asian Conference on Machine Learning*, PMLR 2015; 39: 129-141
- [4] Michel Journee, Yurii Nesterov, Peter Richtarik and Rodolphe Sepulchre; Generalized Power Method for Sparse Principal Component Analysis. *Journal of Machine Learning Research* 2010; 11: 517–553
- [5] Nicolas Gillis, François Glineur; Accelerated Multiplicative Updates and Hierarchical ALS Algorithms for Nonnegative Matrix Factorization. *Neural Computation* 2012; 24 (4): 1085–1105
- [6] <https://www.kaggle.com/code/washingtongold/movie-data-conversion/data>